


# DATA QUEST

SQL for Real-World Thinking

- Aditya Varma



# Preface

## Welcome to Data Quest SQL

This book is written to help you learn SQL as a skill, not as a subject.

SQL is not about writing long programs. It is about understanding data and asking the right questions. It focuses on clarity over complexity. Because of this, learning SQL builds a strong foundation for working with real-world data across different domains.

The chapters in this book are arranged carefully. Each concept is introduced only when it is needed, and every idea is supported by simple queries that explain *why* something works, not just *how* to write it.

You will not just learn how to write queries, but how to think in terms of data. You will learn how information is structured, how relationships are built, and how meaningful insights are extracted.

No prior database knowledge is assumed. You are expected only to read, practice, and stay curious. This book is meant to be read slowly, like a guide. Run the queries. Modify them. Observe what changes.

Take your time. Ask questions. Make mistakes. That is how real understanding is built.

*-Aditya Varma*

## LICENSE & CREDITS

### License:

This work is licensed under the **Creative Commons Attribution–NonCommercial–ShareAlike 4.0 (CC BY-NC-SA 4.0)** license. You may share and adapt this material for non-commercial purposes, provided proper credit is given and derivative works use the same license.

### Credits:

Certain examples, datasets, and reference ideas in this book are inspired by publicly available educational resources and open datasets from platforms such as Kaggle and other learning communities.

All original content, explanations, and structure in this material are created for educational purposes. Proper credit is given wherever applicable. All rights remain with their respective creators.

# Index

 **Understanding Data**

 **Relational Database**

# Understanding DATA!

## What is Data?

Imagine waking up in the morning and checking your phone. The time you see on the screen, the number of notifications, the weather update, even the number of steps you walked yesterday all have one thing in common. They are small pieces of information describing something about the world or about you. These small pieces are what we call **data**.

Data is simply a collection of facts. It does not need to be complicated. Your name is data. Your age is data. The marks you scored in your last exam are data. Even something as simple as whether it is raining or not is also data. These are all observations, captured and recorded in some form.

On their own, these facts do not say much. The number “85” means nothing unless you know it is your exam score. The number “30” could be temperature, age, or price. Data by itself is quiet. It just exists. It waits to be understood.

When we begin to organize and interpret data, it starts becoming useful. It starts telling us something. That transformation is what makes data powerful, and that is exactly what you will learn throughout this course.

## Data and Information: The Subtle Difference

Let's take a small situation. You are given a list of numbers: 85, 90, 78, 92. At first glance, these are just numbers. There is no meaning attached. This is data.

Now someone tells you, “These are the marks of a student in four subjects.” Suddenly, your understanding changes. You might think the student is doing well. You might even calculate the average. The same numbers now carry meaning.

That meaning is called **information**.

Data is raw and unprocessed. Information is what you get after understanding and interpreting that data. This distinction may seem small right now, but it is at the heart of everything you will do with SQL. You will take raw data and turn it into meaningful information.

## **Why Do We Store Data?**

Think about how much you try to remember in a day. Names, tasks, passwords, conversations. Now imagine trying to remember the details of hundreds or thousands of things over weeks or months. It quickly becomes impossible.

This is why we store data. Consider a small grocery shop. Every day, customers come in and buy products. If the shop owner does not record anything, they will have no idea which products sell the most, which ones are rarely bought, or how much profit they are making. Everything becomes guesswork.

Now scale this up to a large company like an online shopping platform. Millions of users, millions of products, and millions of transactions happen every day. Without storing data, such a system would collapse almost instantly.

Data is stored so that it can be remembered, analyzed, and used later. It allows us to look back at what has happened and make better decisions for the future. A company can improve its services. A teacher can track student performance. A doctor can monitor patient history.

In a way, stored data acts like memory, but far more reliable and far more detailed than what a human mind can handle.

## **Data in the Real World**

Once you start noticing it, you will realize that data is everywhere, quietly working behind the scenes. When you use a shopping app, every action is recorded. The products you search for, the items you add to your cart, the things you actually buy, all of this becomes data. This is why the app seems to “know” what you might like next.

When you watch movies on a streaming platform, it remembers what you watched, how long you watched it, and what you skipped. That data is then used to recommend new movies or shows that match your taste.

Even something as simple as booking a ride involves data. Your pickup location, destination, travel time, and fare are all stored. Over time, this data helps improve routes, estimate prices, and manage drivers more efficiently.

Social media platforms take this even further. Every like, comment, share, and scroll is recorded. Even the time you spend looking at a post becomes data. It may feel invisible, but it is constantly being generated.

Without realizing it, you are producing data every single day.

## A Brief Journey Through the History of Data

Data has always been part of human life, even long before computers existed. In ancient times, people recorded information by carving symbols into stone or writing on walls. Kings and rulers kept records of taxes, land, and population using simple tools. These were early forms of data storage.

As societies grew, paper became the main method of recording information. Large registers and books were used to store records of businesses, schools, and governments. While this was an improvement, it had its own limitations. Finding specific information was slow, updating records was difficult, and storing large amounts of data required a lot of physical space.

The real transformation began with computers. Suddenly, data could be stored in digital form. It could be searched quickly, updated easily, and managed efficiently. This led to the creation of databases, which are designed specifically to organize and handle large amounts of data.

Today, we live in a world where data is generated at an enormous scale. Every click, every interaction, every transaction adds to the growing ocean of data. This modern era is often referred to as the age of data.

Some people describe data as the “new oil.” Not because it looks similar, but because of its value. Just like oil, raw data is not very useful until it is processed. But once it is refined and analyzed, it becomes incredibly powerful.

## Different Forms of Data

Not all data looks neat and organized. In fact, most of it is quite messy. Some data is structured, like a table with rows and columns. This is the kind of data you might see in a spreadsheet, where everything is clearly organized. This is also the type of data you will work with most when using SQL.

Other types of data are less organized. For example, messages, images, videos, and audio files do not fit neatly into rows and columns. These are still data, but they require different ways of handling and processing. For now, it is enough to understand that data comes in many forms, but structured data is where your journey with SQL will begin.

## From Data to Databases

Imagine trying to manage student records by writing them on random pieces of paper. Some names are here, some marks are there, and some details are missing altogether. It would be confusing and inefficient. Now imagine organizing all that information into a table where each student has a row, and each detail has its own column. Suddenly, everything becomes clear and easy to manage.

This organized collection of data is called a **database**.

A database allows us to store data in a structured way so that it can be easily accessed, updated, and analysed. It brings order to what would otherwise be chaos.

## **Why This Chapter Matters**

At this point, you might feel that everything here is simple, almost obvious. And that is exactly the point. Before learning how to write SQL queries, you need to understand what you are working with. SQL is not just about writing commands. It is about asking questions to data and getting meaningful answers.

If you understand what data is, why it is stored, and how it is used, the rest of the journey becomes much easier.

## **Closing Thought**

Right now, data may seem like nothing more than numbers, text, and records. Quiet, passive, and uninteresting.

But as you move forward, you will learn how to interact with it. You will learn how to filter it, combine it, and analyse it. And when that happens, something changes. Data stops being silent. It starts telling stories.

## // Why Data Matters

If data were just stored and never used, it would be like writing a diary and never reading it again. The real value of data comes from what we *do* with it.

To understand why data matters, imagine running a small café.

Every day, customers walk in and order coffee, tea, snacks. You notice that sometimes the café is crowded, and sometimes it is empty. You might have a feeling that weekends are busier, or that people prefer cold coffee in the afternoon. But these are just guesses.

Now imagine you start recording everything. What customers ordered, at what time, on which day. After a few weeks, you look at the data and notice clear patterns. Cold coffee sells more between 2 PM and 5 PM. Weekends bring in twice as many customers. One particular snack is rarely ordered.

Suddenly, you are no longer guessing. You are making decisions based on facts. This is why data matters. It replaces assumptions with clarity.

### **Data and Decision Making**

Every important system today runs on decisions, and most of those decisions are powered by data.

Think about a navigation app. When you enter a destination, it suggests the fastest route. It is not guessing. It is using data collected from thousands of users, traffic patterns, road conditions, and time of day.

Banks use data to decide whether to approve a loan. They look at income, past transactions, credit history, and other factors. The decision is based on patterns found in data, not just human judgment.

Even in something as simple as a classroom, a teacher can use data to track student performance. Instead of relying on memory, they can see exactly which topics students struggle with and adjust their teaching.

Data helps answer questions like:

- What is happening?
- Why is it happening?
- What should we do next?

Without data, decisions become guesses. With data, decisions become informed.



## **Data and Analytics**

If storing data is like collecting puzzle pieces, analytics is the process of putting those pieces together to see the full picture.

Analytics is about exploring data to discover patterns, trends, and insights.

Consider a movie streaming platform. Millions of users watch content every day. The platform collects data about what people watch, how long they watch, what they skip, and what they search for.

By analysing this data, the platform can answer questions like:

- Which genres are popular right now?
- Which shows are being watched completely?
- What type of content keeps users engaged longer?

This is how recommendations work. When a platform suggests a movie “you might like,” it is not random. It is based on data analysis.

Businesses use analytics to understand customers, improve products, and increase profits. Governments use it to study population trends. Sports teams use it to analyze player performance.

Analytics turns raw data into useful knowledge.

## **Data and Automation**

Now imagine taking it one step further. Not only are we analyzing data, but we are also using it to make systems act automatically.

This is where automation comes in.

Think about an online shopping app. When a product goes out of stock, the system can automatically detect it and notify the seller. When a user places an order, the system processes payment, updates inventory, and sends a confirmation message without any human involvement.

All of this happens because data is being continuously monitored and used to trigger actions.

Another example is email filtering. Your inbox automatically separates important emails from spam. It learns from data, patterns, and behaviour to make these decisions.

Automation saves time, reduces errors, and allows systems to operate at a scale that humans alone cannot manage.

In simple terms:

- Data tells us what is happening
- Analytics helps us understand it
- Automation helps us act on it

## Types of Data

So far, we have been talking about data as if it is all the same. But in reality, data comes in different forms, and not all of it is neatly organized.

Understanding these types will help you see why SQL is designed the way it is.

- **Structured Data:**

Structured data is the most organized form of data. It is arranged in a clear format, usually in rows and columns, like a table.

Imagine a bank storing customer details. Each customer has:

- a name
- an account number
- a balance

This information fits perfectly into a table. Each row represents a customer, and each column represents a specific type of information. This kind of data is easy to store, search, and analyse. It follows a fixed structure, which makes it ideal for databases and SQL.

Whenever you see data in spreadsheets or tables, you are looking at structured data.

- **Semi-Structured Data:**

Not all data fits perfectly into tables. Sometimes, data has some organization, but not a rigid structure. This is called semi-structured data.

For example, consider a user profile stored in a flexible format. One user might have a phone number and address, while another might only have an email. The structure is not fixed, but there is still some pattern. Formats like JSON and XML are used to represent this kind of data.

It is like having notes written in a consistent style, but not in strict rows and columns. You can still understand it, but it is not as neatly organized as structured data.

- **Unstructured Data**

Now imagine data that has no predefined structure at all. This includes:

- images
- videos
- audio files
- text messages

A photo on your phone is data. A voice recording is data. A long paragraph of text is also data. But you cannot easily put these into a table with rows and columns. They require different methods to store and process.

Interestingly, most of the data in the world today is unstructured. Think about how many photos, videos, and messages are created every second.

## **Real-World Examples of Data Types**

To make this clearer, let's connect these types to real-world systems.

In a bank, most of the data is structured. Account numbers, balances, and transaction records are all neatly organized in tables. This is why SQL is heavily used in banking systems.

On a platform like Instagram, things are more mixed. Usernames and follower counts are structured data. But photos, videos, and captions are unstructured. Comments and messages may fall somewhere in between.

Amazon uses all three types. Product listings and prices are structured. Product descriptions and reviews can be semi-structured. Images and videos of products are unstructured.

These systems combine different types of data to create a complete experience.

## // What is a Database?

Imagine a school office. There are hundreds of students, each with records like name, class, marks, attendance, and contact details. Now imagine all of this information written on random sheets of paper and thrown into a drawer.

Finding one student's record would be a nightmare.

So instead, the school organizes everything neatly. Each student has a proper record. Everything is stored in a structured way so it can be easily found, updated, or reviewed.

That organized collection of data is what we call a **database**. A database is simply a system designed to **store, manage, and organize data efficiently**.

It is not just about storing data. It is about storing it in a way that makes it:

- easy to access
- easy to update
- easy to analyse

Think of a database as a highly disciplined librarian. Nothing is misplaced, everything has a position, and when you ask a question, it knows exactly where to look.

### Why Not Just Use Excel?

At this point, a very natural question comes up. "If databases store tables, why not just use Excel?"

It's a fair question, because Excel *does* look similar to a database at first glance. Rows, columns, and data all sitting neatly in cells. For small tasks, Excel works perfectly fine. You can store marks of a class, track expenses, or manage a small list of items. But problems begin when things grow.

Imagine trying to manage a large online store using Excel. Thousands of products, customers, and orders. Multiple people trying to access and update the file at the same time. The file becomes slow, errors start appearing, and maintaining consistency becomes difficult.

Now imagine a bank trying to store account details in Excel. One mistake, one accidental deletion, and critical data could be lost or corrupted. There is no strong control over who can access or modify what. Databases are built to handle these problems.

They can manage huge amounts of data without slowing down. Multiple users can access the system at the same time without breaking anything. They provide security, consistency, and reliability.

Excel is like a notebook. A database is like a full library system. Both are useful, but they serve very different scales.

## **Database vs File System**

Before databases became common, data was often stored using simple file systems.

A file system is what you use when you store files on your computer. You might have folders like “Documents,” “Photos,” or “Projects,” and inside them, files containing information. This works well for basic storage. But imagine trying to manage complex data using only files.

Suppose you are building a system for a college. You create:

- one file for students
- one file for courses
- one file for results

Now, what happens when you want to find all students enrolled in a specific course? You would have to open multiple files, search manually, and somehow connect the information. This quickly becomes messy and inefficient.

Databases solve this problem by allowing data to be connected and queried easily. Instead of jumping between files, you can ask one clear question and get a precise answer. For example, instead of manually checking multiple files, you can simply ask: “Show me all students enrolled in this course.” And the database will give you the answer instantly.

Another issue with file systems is duplication. The same data might be stored in multiple files, leading to inconsistencies. If one copy is updated and another is not, errors start creeping in. Databases reduce duplication and maintain consistency by organizing data properly and linking related pieces together.

In simple terms:

- File systems store data
- Databases manage data intelligently

## Types of Databases

As data grew in size and complexity, different types of databases are created to handle different needs.

- **Relational Databases (SQL):**

Relational databases are the most common and the most structured type. They store data in tables. Each table has rows and columns. Different tables can be connected using relationships. For example, in an online store: one table stores customers, other stores orders, other stores products.

These tables are connected. A customer places orders. Orders contain products. Everything is linked. This is why they are called **relational** databases.

To work with these databases, we use SQL, which allows us to:

- retrieve data
- insert new data
- update existing data
- delete unwanted data

Relational databases are used in:

- banking systems
- e-commerce platforms
- school management systems

They are reliable, structured, and perfect for handling organized data.

- **NoSQL Databases (Basic Idea):**

As applications grew more complex, especially with social media and large-scale systems, a new type of database became popular: **NoSQL**. Unlike relational databases, NoSQL databases do not rely strictly on tables. They are flexible & can store data in different formats.

For example, instead of forcing all users to have the same structure, a NoSQL database can store different types of data for different users.

This is useful for applications where:

- data structure changes frequently
- large amounts of unstructured data are involved
- flexibility is more important than strict organization

NoSQL databases are commonly used in:

- social media platforms
- real-time applications
- big data systems

For now, you do not need to go deep into this. Just remember: Relational databases are structured and table-based. NoSQL databases are flexible and less structured.

## Examples of Databases

Let's look at some real tools that are used in the industry.

- **MySQL:**

MySQL is one of the most widely used relational databases. It is known for being simple, fast, and reliable. Many websites and applications use MySQL to store their data.

If you have ever used a website that requires login, there is a good chance MySQL is working in the background.

- **PostgreSQL:**

PostgreSQL is another powerful relational database. It is known for being more advanced and feature-rich. It supports complex queries and is often used in systems that require high performance and reliability.

Many modern applications prefer PostgreSQL because of its flexibility and strong capabilities.

- **MongoDB:**

MongoDB is a popular NoSQL database. Instead of storing data in tables, it stores data in a flexible format similar to JSON.

This makes it easier to handle data that does not fit neatly into rows and columns, such as user profiles with varying details.

It is widely used in modern web applications, especially those dealing with large and dynamic data.

## // What is a DBMS?

The DBMS quietly handles many important responsibilities behind the scenes. It allows users to store new data, update existing data, and retrieve information whenever needed. It also ensures that multiple users can work on the same database without interfering with each other.

For example, in a banking system, thousands of people might be checking balances, transferring money, or updating details at the same time. The DBMS makes sure all of this happens smoothly without conflicts or data loss.

It also provides security. Not everyone should have access to everything. A customer can view their account, but only authorized employees can modify certain details. The DBMS controls who can see and change data.

Another important role is maintaining consistency. If data is updated in one place, the DBMS ensures that the change is reflected correctly everywhere it is needed.

In short, the DBMS makes working with data reliable, secure, and efficient.

### Examples of DBMS:

Some common DBMS software used in the real world include:

- MySQL, which is widely used in web applications
- PostgreSQL, known for its advanced features and reliability
- MongoDB, used for flexible and large-scale data

These tools are what developers and companies use to actually work with databases.

## Basic Terminology

Before going further, it is important to become comfortable with a few basic terms. These are the building blocks of everything you will do in SQL.

- **Table:** A table is where data is stored inside a database. It is organized in a structured way, similar to a spreadsheet.  
For example, a “students” table might contain details like names, ages, and marks. Each type of data has its own column, and all related information is grouped together.
- **Row:** A row represents a single entry in a table. If you imagine a table of students, each row would represent one student. It contains all the information related to that specific record.
- **Column:** A column represents a specific type of data within a table. For example, in a student table, columns could be “name,” “age,” and “marks.” Each column stores one kind of information for all rows.
- **Record:** A record is simply another word for a row. It refers to a complete set of information about one item. So, one student’s full details in a table would be called a record.



- **Schema:** A schema defines the structure of a database. It describes how data is organized, including what tables exist, what columns they have, and how they are related to each other.

You can think of a schema as a blueprint. Just like a building plan shows how rooms are arranged, a schema shows how data is structured inside a database.

## **Closing Thought**

At first, these terms may feel simple, almost too basic. But they are the language of databases. As you move forward and start writing SQL queries, you will use these concepts again and again. Tables, rows, columns, and schemas will become second nature.

Right now, you are learning the names of the pieces.

# Relational Database

## Understanding Relationships in Data

Imagine you are managing a college. You have information about students. You also have information about courses. And somewhere, there is information about which student is enrolled in which course. At first, you might think of putting everything into one big table:

- student name
- course name
- marks
- teacher

But very quickly, problems begin to appear.

If one student is enrolled in three courses, their name will be repeated three times. If a course has fifty students, the course name will be repeated fifty times. The data starts becoming repetitive, messy, and harder to maintain.

Now imagine a better approach. You create one table for students, one table for courses, and another table that connects them. Instead of repeating data again and again, you link related information together. This idea of organizing data into separate tables and connecting them using relationships is the foundation of a **relational database**.

## What Does “Relational” Actually Mean?

The word *relational* comes from the idea of **relationships between data**. Instead of storing everything in one place, data is divided into multiple tables, and these tables are connected using common fields.

For example:

- A student has an **ID**
- A course has an **ID**
- An enrolment table connects student IDs with course IDs

Now, instead of repeating full information, the system simply uses these IDs to relate data.

This makes the database:

- more organized
- less repetitive
- easier to update

If a student's name changes, you update it in one place. Everywhere else, the relationship remains intact.

## Why Relational Databases Matter

Relational databases are not just a design choice. They solve real problems that appear when data grows.

Imagine a company managing customer orders. If customer details are stored again and again with every order, even a small mistake can create inconsistencies. One order might have a slightly different spelling of the customer's name. Another might have outdated contact details.

By separating customers and orders into different tables and linking them, the system avoids duplication and keeps data clean. This approach improves:

- accuracy
- consistency
- efficiency

It also makes querying data much easier. Instead of searching through large, repetitive datasets, the system can quickly connect the right pieces of information.

## A Simple Real-World Example

Let's take a familiar system: an online shopping platform. Instead of storing everything in one place, the system is divided into multiple tables:

- Customers (who is buying)
- Products (what is being sold)
- Orders (what was purchased)

Now imagine you want to answer a question: "Which customers bought a specific product?" Without relationships, this would require scanning through large, repeated data. With a relational structure, the system simply connects: customers → orders → products; and gives you the answer efficiently. This is the power of relational thinking.

## The Idea of Breaking Things Down

One of the most important ideas in relational databases is this: **Do not store everything together. Break it into meaningful parts.**

This might feel counterintuitive at first. Many beginners try to put everything into one table because it seems easier. But in reality, breaking data into smaller, connected tables makes systems:

- easier to manage
- easier to scale
- easier to understand

It is similar to organizing your room. Instead of putting everything in one big box, you use different sections for clothes, books, and other items. Finding things becomes faster and more logical.

## **Where You See This in Real Life**

Relational databases are used in almost every system that handles structured data.

- In a bank, customer details, accounts, and transactions are stored in separate tables but connected through relationships.
- In a school, students, teachers, classes, and results are all linked together.
- In a social media platform, users, posts, and comments are stored separately but connected through relationships.

These systems rely on relational concepts to keep data organized and meaningful.

## // Tables and Structure

Imagine you are managing a small school. On your first day, you decide to store student information. You open a notebook and start writing:

“John, 20, Maths, 85”

“Sarah, 19, Science, 90”

“John, 20, English, 78”

At first, it seems fine. But look closer. The same student appears multiple times. Information is repeated. If John’s age needs updating, you now have to change it everywhere. This is where tables come in.

Instead of writing everything randomly, you organize data into a structured format:

| Student ID | Name  | Age |
|------------|-------|-----|
| 1          | John  | 20  |
| 2          | Sarah | 19  |

And then another table:

| Student ID | Subject | Marks |
|------------|---------|-------|
| 1          | Maths   | 85    |
| 1          | English | 78    |
| 2          | Science | 90    |

Now things start making sense. Data is no longer repeated unnecessarily. Each table has a clear purpose. Information is connected, not duplicated. A **table** is simply a structured way of storing data in rows and columns. But more importantly, it is a way of **thinking clearly about data**.

### Table Design

Designing a table is not just about creating columns. It is about deciding *what belongs together* and *what should be separated*. A common mistake beginners make is trying to put everything into one table because it feels easier. But this often leads to messy, repetitive, and confusing data.

Let’s take a real-world example. Imagine an online store. You might be tempted to create one big table like this:

| Customer Name | Product | Price | Order Date |
|---------------|---------|-------|------------|
| Alice         | Phone   | 500   | Jan 1      |
| Alice         | Laptop  | 1000  | Jan 5      |

At first, it works. But soon problems appear:

- Customer data is repeated again and again
- If Alice changes her name or details, you must update multiple rows
- Data becomes harder to maintain

A better design would be:

- **Customers Table**

| ID | Name  |
|----|-------|
| 1  | Alice |

- **Orders Table**

| Order ID | Customer ID | Date  |
|----------|-------------|-------|
| 101      | 1           | Jan 1 |
| 102      | 1           | Jan 5 |

- **Products Table**

| Product ID | Name   | Price |
|------------|--------|-------|
| 1          | Phone  | 500   |
| 2          | Laptop | 1000  |

Now everything is cleaner. Each table has one responsibility. Data is connected instead of repeated.

Good table design is about:

- reducing repetition
- keeping data consistent
- making future changes easier

It may take a little more thinking at the start, but it saves a lot of trouble later.

## // Keys: Giving Identity to Data

Imagine a classroom again. There are many students, and some of them may even share the same name. If you call out “John,” more than one student might respond.

So how do schools solve this? They assign a **roll number** to every student. That number is unique. No two students share it. It becomes their identity inside the system.

This is exactly what keys do in a database. Keys help us uniquely identify data and connect it across tables.

### Primary Key:

A **primary key** is like that roll number. It is a column (or a set of columns) that uniquely identifies each row in a table. No two rows can have the same primary key value, and it cannot be empty.

Imagine a **students** table:

| student_id | name  | age |
|------------|-------|-----|
| 1          | John  | 20  |
| 2          | Sarah | 19  |

Here, **student\_id** is the primary key. Even if there are two students named John, their IDs will always be different. Without a primary key, it becomes very difficult to uniquely identify and manage records. Updates, deletions, and relationships would all become unreliable.

### Candidate Key:

Sometimes, more than one column can uniquely identify a record. For example, in a student table:

- **student\_id** is unique
- email might also be unique

Both of these can act as identifiers. These are called **candidate keys**.

A candidate key is any column that *could* be used as a primary key. Out of all candidate keys, one is chosen as the primary key.

Think of it like this: A person can be identified by passport number, national ID, or email. All are valid, but the system chooses one as the main identifier.

**Foreign Key:**

Now comes the part where tables start talking to each other. A **foreign key** is a column in one table that refers to the primary key in another table.

Let's go back to the school example. You have a **students** table:

| student_id | name  |
|------------|-------|
| 1          | John  |
| 2          | Sarah |

And a marks table:

| student_id | subject | marks |
|------------|---------|-------|
| 1          | Maths   | 85    |
| 1          | English | 78    |
| 2          | Science | 90    |

Here, **student\_id** in the marks table is a **foreign key**. It connects each row in the marks table to a student in the **students** table.

This is how databases avoid repetition. Instead of copying full student details into the marks table, they just use the ID to link the data. Foreign keys create connections. They turn separate tables into a unified system.



## // Relationships: Connecting the Data

Keys help identify data. Relationships define how data is connected. Let's explore the different types of relationships using simple real-world scenarios.

### One-to-One Relationship:

In a one-to-one relationship, one record in a table is linked to exactly one record in another table. Imagine a system where each person has one passport.

| Person | Passport |
|--------|----------|
| John   | P123     |
| Sarah  | P456     |

Each person has only one passport, and each passport belongs to only one person. This type of relationship is not very common but is used when data needs to be separated for clarity or security.

### One-to-Many Relationship:

This is the most common type of relationship. In a one-to-many relationship, one record in a table can be linked to multiple records in another table.

Think about a teacher and students. One teacher teaches many students, but each student is assigned to one teacher. Or consider a customer placing orders. One customer can place multiple orders, but each order belongs to one customer.

This is how it looks conceptually:

- One customer → many orders
- One student → many marks

This type of relationship is everywhere in real-world systems.

### Many-to-Many Relationship:

Now imagine a situation where both sides can have multiple connections. Students and courses are a perfect example. A student can enroll in multiple courses. A course can have multiple students. This creates a many-to-many relationship.

But relational databases cannot directly handle this type of relationship in a single step. So we introduce a third table, often called a "junction table."

For example:

**Students Table**  
| id | name |

**Courses Table**  
| id | course |

**Enrollments Table**  
| student\_id | course\_id |

This third table connects students and courses, allowing multiple connections on both sides.

### **A Small Story to Tie It Together:**

Imagine building a social media platform. You have users, posts, and comments.

Each user has a unique ID (primary key).

Each post is linked to a user (foreign key).

Each comment is linked to both a user and a post.

Now relationships start forming:

- One user → many posts
- One post → many comments
- Many users ↔ many posts (through likes, shares, etc.)

Without keys and relationships, all this data would be scattered and repetitive. With them, everything becomes structured and connected.

## // Constraints: Setting Rules for Data

Imagine you are designing a form for a school admission system. You would never want:

- a student without a name
- two students with the exact same roll number
- missing important information

So you add rules. The form won't accept invalid or incomplete entries.

In databases, these rules are called **constraints**. Constraints ensure that the data being stored follows certain conditions. They protect the database from incorrect or messy data.

### **NOT NULL:**

Some information is too important to be left empty. Imagine a student record without a name. Or a bank account without an account number. That would not make sense. The **NOT NULL** constraint ensures that a column cannot have an empty value.

For example, in a students table, the name field should always have a value. By applying NOT NULL, the database will reject any attempt to insert a record without a name. It is like saying, "This field is required. You cannot skip it."

### **UNIQUE:**

Now imagine assigning roll numbers to students. If two students accidentally get the same roll number, confusion is guaranteed. Records may overlap, updates may affect the wrong person, and the system loses reliability. The **UNIQUE** constraint ensures that all values in a column are different.

For example, an email address in a user system is usually unique. No two users should have the same email. This constraint protects identity. It makes sure that certain fields remain one-of-a-kind.

### **DEFAULT:**

Sometimes, when data is not provided, you still want a value to be filled automatically. Imagine a new user signing up on a website. If they do not specify their account status, the system might automatically set it to "active." This is where the **DEFAULT** constraint is useful.

It assigns a predefined value when no value is given. It is like saying, "If nothing is specified, assume this." This helps keep data consistent and avoids unnecessary empty values.

## **Why Constraints Matter?**

Without constraints, a database becomes unreliable very quickly. You might end up with:

- missing critical information
- duplicate entries
- inconsistent data

Constraints act like guards at the gate. They ensure that only valid and meaningful data enters the system.

## // Normalization: Keeping Data Clean

Now let's talk about something slightly bigger than constraints. Imagine you are storing customer and order data in one table:

| Customer Name | Product | Price |
|---------------|---------|-------|
| Alice         | Phone   | 500   |
| Alice         | Laptop  | 1000  |

At first, it looks fine. But notice something:

- "Alice" is repeated multiple times
- If her name changes, you must update multiple rows
- If you miss one update, data becomes inconsistent

This repetition is called **data duplication**, and it creates problems. Normalization is the process of organizing data to **reduce duplication and improve consistency**. It is like cleaning and organizing a messy room so that everything has its own place.

### Why Avoid Duplicate Data?

Duplicate data might seem harmless at first, but it leads to serious issues over time. If the same information exists in multiple places:

- updating becomes difficult
- mistakes become more likely
- storage is wasted
- data becomes unreliable

By reducing duplication, we make the system cleaner and easier to manage.

### First Normal Form (1NF):

The first step in normalization is called **First Normal Form (1NF)**. The idea is simple: Each column should contain **atomic values**, meaning no multiple values in a single field.

Imagine a table like this:

| Name | Subjects       |
|------|----------------|
| John | Maths, English |

Here, the "Subjects" column contains multiple values. This makes it harder to query and manage.

In 1NF, we break this into:

| Name | Subject |
|------|---------|
| John | Maths   |
| John | English |

Now each cell contains a single value. The data becomes easier to work with. 1NF is about making sure data is stored in a clean, simple structure.

### **Second Normal Form (2NF):**

Once data is in 1NF, the next step is **Second Normal Form (2NF)**. This focuses on removing **partial dependency**, which simply means that data in a table should depend on the entire key, not just part of it.

Let's simplify this with an example. Imagine a table:

| Student ID | Course | Student Name |
|------------|--------|--------------|
|------------|--------|--------------|

Here:

- Student ID + Course together identify each row
- But "Student Name" depends only on Student ID, not on the course

This creates unnecessary repetition of the student's name.

To fix this, we split the table:

**Students Table**  
| Student ID | Name |

**Enrollments Table**  
| Student ID | Course |

Now, each piece of data lives where it belongs. 2NF is about separating data so that each table has a clear purpose and no unnecessary repetition.

### **Closing Thought:**

Constraints and normalization may seem like rules and restrictions, but they actually make your life easier. They prevent errors before they happen. They keep data clean and meaningful. They make systems easier to manage as they grow.

At this stage, you are learning how to design data properly. Very soon, you will start using SQL to interact with that data. And when the structure is clean, everything you do with SQL becomes smoother, faster, and far more powerful.

## // Relational Database Systems in the Real World

So far, we've been talking about databases in a general sense. But in reality, you never interact with an "abstract database." You always use a **database system**, which is actual software installed on a computer.

Think of it like this: A database is the idea of storing organized data. A database system is the actual tool that makes it possible. These systems are used everywhere, from small apps to massive global platforms.

### Database vs SQL (Important Distinction):

Before going further, it is very important to clear one common confusion. A **database** is where data is stored. **SQL** is the language used to interact with that data.

It's similar to this:

- A library stores books
- Language is what you use to ask the librarian for a book

You don't "store data in SQL." You use SQL to talk to a database system.

For example:

- PostgreSQL stores the data
- SQL is used to query that data

Once students understand this difference, everything becomes much clearer.

### Popular Relational Database Systems

Let's explore some of the most widely used relational databases in the world. These are not just academic tools. These are the engines running real applications you use every day.



- **MySQL:**

MySQL is one of the most popular relational databases, especially in web development. It is known for being simple to use, fast, and reliable. Many websites and applications use MySQL as their backend database.

Imagine a basic e-commerce website. When a user logs in, views products, or places an order, all that data is stored and managed using a system like MySQL. It became extremely popular because it is open-source and works well with many programming languages.

Real-world usage includes:

- small to medium websites
- blogging platforms
- basic web applications

If you have ever used a login system on a website, there is a high chance MySQL was involved somewhere behind the scenes.

- **PostgreSQL:**

PostgreSQL, often called “Postgres,” is a more advanced and powerful relational database. It is designed for systems that require complex operations, large-scale data handling, and strong reliability. While MySQL focuses on simplicity, PostgreSQL focuses on **capability and precision**.

For example, a financial system that tracks transactions must be extremely accurate. Even a small error can cause serious problems. PostgreSQL is often chosen in such cases because of its strong support for data integrity and advanced features.

Real-world usage includes:

- banking systems
- large-scale applications
- data-heavy platforms

Many modern companies prefer PostgreSQL because it can handle both simple and complex tasks very well.

- **SQLite:**

SQLite is quite different from the others. It is a lightweight, file-based database. This means it does not run as a separate server. Instead, it exists as a simple file on your system. Because of this, it is extremely easy to use and requires almost no setup.

Imagine building a mobile app. You need to store user data locally on the device. Running a full database server would be unnecessary. Instead, SQLite can handle everything inside a small file. Many apps on your phone use SQLite without you even realizing it.

Real-world usage includes:

- mobile applications
- small desktop applications
- embedded systems

- **Oracle Database:**

Oracle is a powerful, enterprise-level database system. It is used by large organizations that deal with massive amounts of data and require high performance, security, and reliability.

Unlike MySQL or PostgreSQL, Oracle is not typically used for small projects. It is designed for large-scale systems like corporate applications, government databases, and major financial institutions. Imagine a multinational company managing millions of transactions every day across different countries. Oracle is built for that level of complexity.

Real-world usage includes:

- large enterprises
- government systems
- high-security environments

- **Microsoft SQL Server:**

SQL Server is widely used in companies that rely on Microsoft technologies. It integrates well with tools like .NET and Azure, making it popular in enterprise environments.

You will often find it in:

- corporate systems
- internal business applications
- enterprise dashboards

- **Amazon Aurora:**

Aurora is a cloud-based database built by Amazon. It is compatible with MySQL and PostgreSQL but designed to handle large-scale cloud workloads efficiently.

Used in:

- cloud applications
- scalable web systems

- **Google Cloud SQL:**

This is not a new database engine, but a managed service that runs MySQL, PostgreSQL, or SQL Server in the cloud. It removes the need to manage infrastructure manually.

### **Putting It All Together:**

At this point, the landscape might feel large, but there is a simple way to think about it.

- Small apps → SQLite
- Web apps → MySQL / PostgreSQL
- Growing systems → PostgreSQL / MariaDB
- Enterprises → Oracle / SQL Server / Db2
- Cloud systems → Aurora / Cloud SQL
- Distributed systems → CockroachDB

Different tools for different needs.



# SQL

At this point in your journey, you already understand that data lives inside databases, neatly organized into tables with rows and columns. But having data stored somewhere is only half the story. The real question is: *how do we actually interact with it?*

This is where SQL comes in.

SQL, which stands for **Structured Query Language**, is the language we use to communicate with a database. It acts as a bridge between you and the stored data. Without SQL, the database would just sit there silently, holding information but offering no easy way to access or modify it.

Imagine walking into a massive library with millions of books but no catalog system, no librarian, and no way to search. You would be surrounded by information, yet unable to use it effectively. SQL solves this exact problem for databases. It allows you to ask precise questions and get exact answers.

SQL is a language used to **communicate with a database**. It allows you to:

- ask questions
- retrieve data
- add new data
- update existing data
- delete unwanted data

If a database is where data lives, then SQL is how you talk to it.

When you write SQL, you are not writing long programs. You are forming requests. These requests can be as simple as “show me all records” or as specific as “show me customers who bought a product in the last 7 days and spent more than a certain amount.”

## Understanding “Structured Query Language”

The name itself tells you a lot about what SQL does.

“Structured” refers to the fact that SQL works with structured data, meaning data that is organized in tables. Unlike random text or images, this kind of data follows a clear format, which makes it easier to search and analyse.

“Query” means asking a question. In SQL, a query is simply a request for data. Every time you use SQL, you are essentially asking the database something.

“Language” means it has its own syntax and rules, just like any other programming language. But unlike most programming languages, SQL is designed to be closer to human thinking. Many SQL statements read almost like plain English, which makes them easier to understand and remember.

When you write a query like: `SELECT name FROM students WHERE age > 18;`

You are simply saying: “Give me the names of students whose age is greater than 18.”

You are not telling the database how to search through rows or how to compare values internally. The database system figures that out for you.

This brings us to one of the most important ideas in SQL.

## **The Declarative Nature of SQL (WHAT vs HOW)**

One of the most important ideas in SQL, and one that often feels different at first, is that SQL is a **declarative language**.

In most programming languages like Java or Python, you are required to tell the computer exactly how to perform a task. You break down the process step by step. For example, if you want to find all students older than 18, you would typically loop through a list, check each value, and store the results.

In SQL, you don’t do that. Instead, you simply state what you want, and the database system figures out how to get it.

For example, when you write a query to find students above a certain age, you are not telling the database how to scan rows, how to compare values, or how to optimize the search. You are only describing the result you want. The database engine handles the entire process internally. This is a powerful idea because it allows you to focus on the problem rather than the steps. It also allows the database system to optimize performance in ways that would be difficult to do manually.

It is **declarative**, which means you only describe *what* you want, not *how* to get it. You say: “Give me this data.”

The database decides:

- how to search
- how to filter
- how to optimize the process

This makes SQL powerful and efficient, especially when dealing with large amounts of data.

A good way to think about this is ordering food in a restaurant. You don’t go into the kitchen and explain how to cook the dish. You simply place your order. The kitchen handles the rest. SQL works in the same way.

## Examples of SQL in Action

To understand SQL better, it helps to see how naturally it expresses ideas. If you want to see all the data in a table, you write a query that essentially says, “show me everything from this table.” The structure of the language reflects this intention very clearly. If you want only certain columns, you specify exactly what you need. If you want to filter data based on a condition, you describe that condition directly.

For example, if you are working with student data and want to see only students above a certain age, your query expresses that requirement directly. There is no need to manually search through each record. The database does that for you.

Let’s look at how natural SQL feels once you understand it.

If you want to see all data in a table: `SELECT * FROM students;`

This means: “Show me everything from the students table.”

If you want specific columns: `SELECT name, age FROM students;`

This means: “Show me only the name and age.”

If you want to filter data: `SELECT name FROM students WHERE age > 18;`

This means: “Show me names of students older than 18.”

Notice how close this is to normal English. That is one of the reasons SQL is considered easy to read and write compared to many programming languages. What makes SQL interesting is how closely it mirrors human thinking. You are not writing instructions for a machine step by step. You are expressing a requirement in a structured form.

## Types of Operations in SQL

Although SQL feels simple on the surface, it is capable of handling many different types of operations. Sometimes you will use SQL to retrieve data. This is the most common use case, where you ask questions and get results. Other times, you will use it to modify data. You might add new records, update existing ones, or remove data that is no longer needed. You can also use SQL to define the structure of your database itself. This includes creating tables, modifying them, and setting rules for how data should be stored.

So, while SQL may look like a simple query language, it actually covers a wide range of actions, all centered around managing and working with data.

You don’t need to memorize everything right now, but it’s helpful to see the big picture. Some commands are used to **retrieve data**, like `SELECT`. Some are used to **modify data**, like:

- `INSERT` (add data)
- `UPDATE` (change data)
- `DELETE` (remove data)

Others are used to **define structure**, like:

- **CREATE TABLE**
- **ALTER TABLE**

As you move forward, you will explore each of these in detail.

## **SQL in the Real World**

SQL is not just something used in classrooms or small projects. It is deeply embedded in real-world systems.

When you log into a website, SQL is often used to check your credentials. When you browse products on an online store, SQL queries are used to fetch product details. When a company analyzes customer behaviour, SQL is used to retrieve and process the necessary data.

Even applications that seem very modern and complex still rely on SQL at their core when it comes to handling structured data. This is why SQL has remained relevant for decades. It solves a very fundamental problem: how to work with organized data efficiently.

Imagine you are working in a company that has thousands of customers. Your manager asks: “Can you give me a list of customers who placed orders in the last 7 days?”

Without SQL, this would be extremely difficult. You would have to manually search through records, match data, and compile results. With SQL, it becomes a single query. You ask the database the question, and it gives you the answer.

This is why SQL is used everywhere:

- in banking systems
- in e-commerce platforms
- in social media
- in analytics and reporting

Wherever there is structured data, SQL is usually there.

## **A Few Interesting Facts About SQL**

SQL has been around since the 1970s, yet it continues to be one of the most widely used technologies in the world. Few technologies last that long without becoming outdated.

Another interesting fact is that while there are many database systems, the core ideas of SQL remain largely the same across all of them. This means that once you learn SQL, you are not tied to a single tool. You can apply your knowledge in different environments with only small adjustments.

It is also worth noting that SQL is often used by people who are not traditional programmers. Analysts, data scientists, and even business professionals use SQL because of its clarity and practical value.

## // SQL vs Programming Languages

By now, you might be wondering something natural: “If SQL is a language... is it like Java or Python?”. The short answer is: **not really**.

SQL is a language, yes. But it behaves very differently from traditional programming languages. If you try to treat SQL like Java or Python, it will feel confusing. But once you understand its nature, it becomes surprisingly simple.

### **Different Purpose, Different Thinking**

Programming languages like Java or Python are designed to build complete applications. You use them to define logic, control flow, create functions, and manage how a program behaves step by step.

SQL, on the other hand, is not designed to build applications. It is designed for one main purpose:

### **Working with data**

You don't use SQL to create full software systems. You use it to interact with the data inside those systems.

Think of it like this:

- Java/Python → building the entire machine
- SQL → accessing and managing the machine's data

Both are important, but they play very different roles.

### **SQL is Not Step-by-Step Programming**

In programming languages, you often write instructions in sequence. For example, if you want to find students above 18 in Python, you would:

- loop through each student
- check their age
- store matching results

You are controlling the entire process step by step.

In SQL, you skip all of that. You simply say: “Give me students whose age is greater than 18.”. And the database handles everything internally.

This is why SQL feels simpler. You are not managing the process, you are defining the result.

## **No Loops (Mostly)**

One of the biggest differences students notice is this: SQL does not use loops in the way programming languages do.

There is no need to say:

- “for each row, do this”
- “repeat this process”

Why? Because SQL is built to work with entire datasets at once. When you write a query, you are not thinking about one row at a time. You are thinking about *all rows together*.

For example, when you ask for students above 18, SQL checks all rows internally. You don’t need to write a loop for it.

It’s like asking: “Show me all red cars in the parking lot.” You don’t go car by car. You let the system filter everything at once.

## **Set-Based Thinking (The Real Shift)**

This brings us to an important concept. Programming languages often use **step-by-step thinking**. SQL uses **set-based thinking**.

A “set” simply means a collection of data. Instead of thinking: “Take one row, process it, then move to the next”. You think: “Work with the entire group of rows together”

This is why SQL is so powerful when dealing with large amounts of data. It can handle thousands or even millions of rows in a single query without you writing complex logic.

## **Focus on Data, Not Logic**

In programming, a lot of your effort goes into:

- writing logic
- controlling flow
- handling conditions

In SQL, your focus shifts completely.

You think about:

- what data you need
- where it is stored
- how it is related

You are not building logic-heavy programs. You are asking meaningful questions. This is why SQL is widely used not only by developers, but also by analysts and business professionals.

## **A Simple Comparison**

To make it clearer, let's put it side by side in a simple way:

- In Python, you might write a loop to filter data
- In SQL, you write one query and get the result
- In Java, you define classes, methods, and control flow
- In SQL, you define what data you want
- Programming languages tell the computer *how to do something*
- SQL tells the database *what you want*

## **A Small Story**

Imagine you are working in a company and your manager asks: "Can you find all customers who made purchases this month?"

If you were using a programming language, you might:

- load all data
- loop through records
- apply conditions
- collect results

With SQL, you simply write one query. You ask the question directly, and the database gives you the answer. Less effort, more clarity.

## **Closing Thought**

SQL is not here to replace programming languages. It works alongside them. Programming languages build systems. SQL powers those systems with data.

Once you understand this difference, SQL stops feeling strange. Instead, it starts feeling... efficient. You are no longer telling the system what to do step by step. You are simply asking the right question and letting the system do the heavy lifting.

## // **SQL Tools Setup: Getting Your Environment Ready**